

Creation of resource_group myResourceGroup in eastus location (Resource group - logical group in which Azure resources are deployed and managed).

```
az group create --name myResourceGroup --location eastus
```

Enable_cluster_monitoring:

1. Check if Microsoft.OperationsManagement and Microsoft.OperationalInsights are registered on subscription

```
az provider show -n Microsoft.OperationsManagement -o table  
az provider show -n Microsoft.OperationalInsights -o table
```

2. If not registered, register them:

```
az provider register --namespace  
Microsoft.OperationsManagement  
az provider register --namespace Microsoft.OperationalInsights
```

3. Instantiate AKS Cluster

```
az aks create --resource-group myResourceGroup --name  
myAKSCluster --node-count 1 --enable-addons monitoring --  
generate-ssh-keys
```

i. `--enable-addons monitoring` - to enable Azure monitor for containers (<https://docs.microsoft.com/en-us/azure/azure-monitor/containers/container-insights-overview>)

Connect_to_the_cluster:

1. Configure **kubectl** to connect to kubernetes cluster
1. Download credentials and to configure **kubectl** (Updates `~/.kube/config` (configuration file used by kubectl))

```
az aks get-credentials --resource-group myResourceGroup --name  
myAKSCluster
```

2. Verify connection to the cluster

```
kubectl get nodes
```

3. Run application:

i. `vi Dockerfile`

```
FROM php:5-apache  
COPY index.php /var/www/html/index.php  
RUN chmod a+rx index.php
```

ii. `vi index.php`

```
<?php
    $x = 0.0001;
    for ($i = 0; $i <= 1000000; $i++) {
        $x += sqrt($x);
    }
    echo "OK!";
?>
```

iii. `vi php-apache.yaml`

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: php-apache
spec:
  selector:
    matchLabels:
      run: php-apache
  replicas: 1
  template:
    metadata:
      labels:
        run: php-apache
    spec:
      containers:
        - name: php-apache
          image: k8s.gcr.io/hpa-example
          ports:
            - containerPort: 80
          resources:
            limits:
              cpu: 500m
            requests:
              cpu: 200m
---
apiVersion: v1
kind: Service
metadata:
  name: php-apache
  labels:
    run: php-apache
spec:
  ports:
    - port: 80
  selector:
    run: php-apache
```

iv. Deploy the application:

```
kubectl apply -f php-apache.yaml
```

4. Creation of Horizontal Pod Autoscaler

```
kubectl autoscale deployment php-apache --cpu-percent=50 --min=1  
--max=10
```

5. Check current status of autoscaler:

```
kubectl get hpa
```

6. Update existing AKS cluster to enable cluster autoscaler

```
az aks update \  
  --resource-group myResourceGroup \  
  --name myAKSCluster \  
  --enable-cluster-autoscaler \  
  --min-count 1 \  
  --max-count 3
```

i. It takes a few minutes to update cluster and configure cluster autoscaler settings

7. Increase Load

```
kubectl run -i --tty load-generator --rm --image=busybox --  
restart=Never -- /bin/sh -c "while sleep 0.01; do wget -q -O-  
http://php-apache; done"
```

i. Within a minute or so, we should see higher CPU load by executing the following:

```
kubectl get hpa  
kubectl get deployment php-apache  
kubectl get nodes
```

8. Stop load

```
kubectl get hpa  
kubectl get deployment php-apache  
kubectl get nodes
```

To Scale Up Faster.:

1. Update Horizontal Pod Autoscaler with more Pods:

```
kubectl edit hpa php-apache
```

i. Change the following entries:

```
spec:
```

```
maxReplicas: 20
...
targetCPUUtilizationPercentage: 20
```

2. Update Cluster Autoscaler with More Nodes:

```
az aks update --resource-group myResourceGroup --name
myAKSCluster --update-cluster-autoscaler --max-count 4
```

3. Increase Load

```
kubectl run -i --tty load-generator --rm --image=busybox --
restart=Never -- /bin/sh -c "while sleep 0.01; do wget -q -O-
http://php-apache; done"
```

Delete the setup:

```
az group delete --name myResourceGroup --yes --no-wait
```